

# High-Performance Privacy-Preserving Blockchain via GPU-Accelerated Fully Homomorphic Encryption

Anonymous Authors

Anonymous Institute

**Abstract.** Data privacy is essential for safety-critical blockchain applications such as digital payment. A promising approach for preserving privacy is to encrypt transaction data using homomorphic encryption (HE) and prove the correctness of transaction execution through non-interactive zero-knowledge proofs. However, previous works on this approach are restricted by the use of HE schemes that only support addition or multiplication, making it challenging to implement business logic involving both types of arithmetic operations. In addition, previous works suffer from poor performance caused by the costly HE computation, hindering their adoption for real-world applications.

We present GAFE<sup>1</sup>, a high-performance privacy-preserving blockchain that is co-designed with GPU-accelerated fully homomorphic encryption (FHE). GAFE encrypts transaction data with FHE, which allows both addition and multiplication on ciphertexts. For high performance, GAFE leverages parallel execution on GPUs to accelerate FHE computations. To ensure result correctness, GAFE generates lightweight NIZKPs that incur low overhead. Evaluations show that GAFE is highly performant, achieving a  $3.1\times$  increase in throughput (258 transactions per second) and a 37% reduction in latency (1.61 seconds), surpassing the baseline.

**Keywords:** Blockchain · Privacy Preserving · GPU Acceleration · Fully Homomorphic Encryption

## 1 Introduction

Blockchain has been an appealing solution for various industries due to its exceptional characteristics, such as data transparency. However, notable blockchains like Hyperledger Fabric [2] face a serious privacy issue as they process and store transaction data in plaintext, exposing sensitive information to anyone with access to the blockchain. This privacy issue impedes the widespread adoption of blockchain, especially in safety-critical applications such as finance [23], where the privacy of transaction data is crucial.

A promising approach to addressing the privacy issue is using homomorphic encryption (HE) [1] and non-interactive zero-knowledge proofs (NIZKPs) [26].

---

<sup>1</sup> GAFE stands for GPU-Accelerated Fully Homomorphic Encryption Blockchain.

HE enables arithmetic computation to be performed directly on ciphertexts. NIZKPs allow a prover to prove that a statement is true without revealing the involving private inputs. In this approach, transaction data is encrypted using an HE scheme, and transactions are executed directly on ciphertexts. Moreover, NIZKPs are generated to prove the correctness of transaction execution, without disclosing the plaintexts of the execution results.

However, previous works on this approach face two non-trivial challenges. Firstly, they have serious limitations in expressing complex business logic due to their use of partly homomorphic encryption (PHE) to encrypt transaction data. PHE only supports a single type of arithmetic operations, either addition or multiplication. Thus, these works are not suitable for applications that involve both types of arithmetic operations. Secondly, previous works suffer from low performance due to the intensive computational overhead of HE. For instance, the notable privacy-preserving blockchain ZeeStar [25] achieves a long latency of tens of seconds. As a result, these works are unable to meet the high-performance requirements expected in many blockchain applications, such as digital payment.

In this paper, our key insight to address the challenges is that, *we can co-design blockchain with GPU-accelerated fully homomorphic encryption (FHE)*. FHE supports both addition and multiplication on ciphertexts, thereby facilitating the implementation of complex business logic involving both types of arithmetic operations. To enhance performance, we can concurrently execute multiple transactions and parallelize a substantial portion of the arithmetic calculations involved in FHE schemes [28]. These parallel computations are well-suited for GPUs, which offer superior parallel processing capabilities and are widely available in modern commodity machines. Hence, employing GPU-accelerated FHE is both beneficial and feasible for blockchains, as it enables high performance, privacy preservation, and the implementation of complex business logic.

Nonetheless, the trivial combination of blockchain and FHE leads to the problem of inconsistent ciphertext results. In particular, when given identical inputs, different executor nodes may generate inconsistent ciphertexts for the same plaintext result. This inconsistency arises from the inclusion of random noise in FHE ciphertexts for security purposes [27].

To solve the problem of inconsistency, our second insight is *to integrate lightweight NIZKPs with the trusted execution mechanism of the execute-order-validation (EOV) blockchain workflow* [2]. This mechanism first executes a transaction on multiple executors and deems it correct only if a majority of executors produce consistent results. Consequently, we can generate lightweight NIZKPs to prove that a majority of ciphertext results are decrypted to the same plaintexts. The generation of NIZKPs incurs minimal overhead as it does not require costly FHE arithmetic computations.

These insights lead to GAFE, a high-performance privacy-preserving blockchain. GAFE carries a *GPU-accelerated transaction execution workflow*, which comprises four phases. First, the client encrypts transaction data using FHE and sends the encrypted transaction to all executors. Second, each executor leverages a GPU to execute FHE computation for multiple transactions in parallel,

System	FHE Support	GPU Acceleration	High Performance
❖ ZeeStar [25]	x	x	x
❖ Ekiden [9]	x	x	✓
❖ Hawk [18]	x	x	✓
❖ Arbitrum [16]	x	x	✓
◆ Zcash [14]	x	x	x
◆ Monero [20]	x	x	x
◆ FabZK [17]	x	x	✓
◆ Zether [7]	x	x	x
◆ RFPB [29]	x	x	✓
◆ GAFE	✓	✓	✓

Table 1: Comparison of GAFE and related privacy-preserving blockchains. ❖ and ◆ represent general-purpose and specific-purpose blockchains, respectively.

while the client generates NIZKPs to prove the correctness of execution. Third, the orderers run a consensus protocol to determine the transaction order within each block. Finally, the executors validate the transactions in the predetermined order. We built GAFE on the codebase of Hyperledger Fabric and conducted performance evaluations. The results show that GAFE achieves high performance, with an effective throughput of 258 transactions per second and a commit latency of 1.61 seconds. Compared to the baseline system, GAFE exhibited a notable  $3.1\times$  increase in throughput and a significant 37% reduction in latency.

In summary, we make the following contributions: (1) We propose a GPU-accelerated transaction execution workflow that integrates GPU-accelerated FHE into blockchain and ensures the execution correctness through lightweight NIZKPs. (2) We implement GAFE, a high-performance privacy-preserving blockchain that incorporates the aforementioned workflow. (3) We conduct evaluations on GAFE, demonstrating its effectiveness and high performance.

## 2 Related work

### 2.1 Privacy-Preserving Blockchain

We categorize previous works on privacy-preserving blockchains into two groups, as shown in Table 1.

**General-purpose blockchains.** ZeeStar [25] uses ElGamal encryption [19], which only supports addition on ciphertexts. While ZeeStar extends addition to emulate multiplication, this extension is inefficient and not applicable to ciphertexts encrypted by different keys. Ekiden [9] requires hardware that supports trusted execution environments (TEEs) to execute transactions involving private data off the blockchain. However, such TEE-enabled hardware may not always be available for commodity desktop machines. Hawk [18] and Arbitrum [16] delegate transaction execution to trusted managers, but this introduces a substantial attack surface for malicious managers to expose private inputs.

**Specific-purpose blockchains.** To ensure the privacy of digital payment information like the payment amount, Zcash [14] employs NIZKPs, while Monero [20] utilizes ring signatures and stealth addresses. However, their low throughput and

long latency have hindered their broader adoption [24,7]. FabZK [17] combines NIZKPs with a specialized tabular ledger data structure to preserve data privacy, but this data structure restricts FabZK to performing only HE addition. Zether [7] and RFPB [29]<sup>2</sup> adopt PHE schemes, which means they share a limitation with ZeeStar: they only support on-ciphertext addition.

## 2.2 GPU-Accelerated Fully Homomorphic Encryption

Many works have been proposed to leverage GPU acceleration for FHE in order to unlock the potential of FHE in practical applications. HE-Booster [28] accelerates polynomial arithmetic computation by mapping five common phases of typical FHE schemes to the GPU parallel architecture. HE-Booster also introduces thread-level and data-level parallelism to enable acceleration on single-GPU and multi-GPU setups, respectively. Ozcan et al. [22] presents a library that makes efficient use of the GPU memory hierarchy and minimizes the number of GPU kernel function calls. Yang et al. [30] provides GPU implementations of three notable FHE schemes (BGV [6], BFV [5,12], and CKKS [10]) along with various theoretical and engineering optimizations, including a hybrid key-switching technique and several kernel fusing strategies.

## 3 Overview

### 3.1 System Model

GAFE comprises three types of participants: *client*, *executor*, and *orderer*. Executors and orderers are commonly referred to as *nodes*. GAFE uses permissioned settings, where all participants are organized into distinct *organizations*. Each organization runs multiple executors and orderers, and also possesses a set of clients. We built GAFE on top of Hyperledger Fabric [2] and developed a prototype system implementing the business logic of digital payment. Specifically, the role of each participant type is described as follows:

**Client.** Clients encrypt transaction data and submit the encrypted transactions to the nodes for execution and commitment. Each client is identified by a unique fixed-length string *id* and possesses a randomly generated a public-private key set that is used for performing FHE operations. Additionally, each client is associated with a floating-point number *bal*, which represents the balance owned by the client and is encrypted using the FHE scheme.

**Executor.** Each executor is responsible for three main tasks: (1) executing encrypted transactions, (2) validating transaction results, and (3) maintaining the latest local copy of blockchain and state database. Specifically, each executor is equipped with a GPU and employs GPU acceleration for FHE computation during transaction execution. We provide a detailed discussion of GAFE’s transaction execution workflow in §4.2.

---

<sup>2</sup> We refer to the system proposed in [29] as RFPB for convenience.




Fig. 1: GAFE’s GPU-accelerated transaction execution workflow. Each executor utilizes multiple Streaming Multiprocessors (SM) of a GPU to concurrently execute FHE computation for multiple transactions.

**Orderer.** GAFE orderers determines the order of transactions within each block. GAFE orderers leverage the orderer implementation of Hyperledger Fabric, therefore inheriting the mature security guarantees. Specifically, GAFE orderers run an instance of the BFT-SMaRT [4] protocol, a fast Byzantine-Fault-Tolerant consensus protocol. This eliminates the performance bottleneck associated with time-consuming consensus protocols like Proof-of-Work [21].

### 3.2 Threat Model

GAFE adopts the Byzantine failure model [3,8], which tolerates up to  $N$  malicious orderers out of a total of  $3N + 1$  orderers. Participants within an organization trust each other, but not participants from other organizations. We make standard assumptions on FHE and NIZKP.

### 3.3 GAFE’s Workflow Overview

GAFE’s workflow consists of two sub-workflows: *client key generation* (§4.1) and *GPU-accelerated transaction execution* (§4.2). Specifically, the online transaction execution workflow is divided into four phases, as illustrated in Figure 1.

Phase 1: Construction. The client constructs a transaction  $tx$  by encrypting the transaction data (e.g., payment amount  $val$ ) using the client’s encryption key. The client also generates an NIZKP  $\pi_{bal \geq val}$  to demonstrate that the client’s balance  $bal$  is greater than or equal to  $val$ . Finally, the client submits the encrypted transaction and  $\pi_{bal \geq val}$  to all executors for execution.

**Phase 2: Execution.** The executor verifies  $\pi_{bal \geq val}$  and proceeds with the remaining procedure only if  $\pi_{bal \geq val}$  is valid. Next, the executor buffers  $tx$  along with other transactions received within a specific time frame. These buffered transactions are delivered to a GPU for performing concurrent FHE computations, and their execution results are returned to the corresponding clients. Upon receiving the results from all executors, the client generates correctness NIZKPs to prove the correctness of transaction execution. Finally, the client sends both the consistent result and the correctness NIZKPs to the orderers.

**Phase 3: Ordering.** All orderers run a BFT consensus protocol to determine the order of transactions within each block. Once the order is determined, the orderers deliver the block to the executors for validation.

**Phase 4: Validation.** On receiving a block from the orderers, the executor sequentially validates each transaction within the block following the determined order. The executor commits a transaction only if the transaction has valid correctness NIZKPs and has no write conflict with previous committed transactions within the same block. Otherwise, the executor aborts the transaction.

## 4 Workflow Description

### 4.1 Client Key Generation

In GAFE, each client possesses a unique public-private key set for FHE operations. This key set consists of (1) an encryption key  $pk$ , (2) a decryption key  $sk$ , and (3) an evaluation key  $ek$  used for on-ciphertext arithmetic evaluation. While the encryption key  $pk$  and the evaluation key  $ek$  are accessible to all participants in the GAFE network, the decryption key  $sk$  must remain private to the client.

Upon the creation of a client, GAFE runs an instance of the key generation algorithm associated with the chosen FHE scheme (e.g., CKKS [10]). This algorithm initially generates the encryption key  $sk$ , and then derives the decryption key  $pk$  and the evaluation key  $ek$  from  $sk$ . Additionally, GAFE generates a unique fixed-length string  $id$  based on  $sk$  to serve as the client’s unique identifier.

### 4.2 GPU-Accelerated Transaction Execution

GAFE’s *GPU-accelerated transaction execution* workflow co-designs the execute-order-validate workflow of Hyperledger Fabric with GPU-accelerated FHE scheme. This workflow consists of four consecutive phases, as outlined in Figure 1 and Algorithm 1.

**Phase 1: Construction.** For privacy preservation, GAFE processes and stores transaction data in the form of ciphertext. Consequently, the client is required to construct an encrypted transaction and submit it to GAFE nodes for execution.

**Phase 1.1: Encrypting transaction data.** The client encrypts transaction data using the encryption keys of the clients involved. In the context of digital payment, a transaction involves a sender client  $c_s$  and a receiver client  $c_r$ . Therefore,  $c_s$  is responsible for encrypting the payment amount  $val$  into two ciphertexts  $val_s$

**Algorithm 1:** Transaction execution workflow of the client.

---

```

input: Plaintext transaction data  $val$ 
// Phase 1: Construction
1  $pk_s \leftarrow \text{GetEncryptionKey}(id_s); val_s \leftarrow \text{Encrypt}(val, pk_s);$ 
2  $pk_r \leftarrow \text{GetEncryptionKey}(id_r); val_r \leftarrow \text{Encrypt}(val, pk_r);$ 
3  $bal_s \leftarrow \text{GetBalance}(id_s);$ 
4  $\pi_{bal \geq val} \leftarrow \text{GenerateNIZKP}(bal_s \geq val_s, sk_s);$ 
5  $tx \leftarrow \{val_s, val_r, \pi_{bal \geq val}\};$ 
// Phase 2: Execution
6  $results \leftarrow \emptyset;$ 
7 For every executor e; do in parallel
8    $bal'_s, bal'_r \leftarrow \text{SendForExecution}(e, tx);$ 
9    $results \leftarrow results \cup \{bal'_s, bal'_r\};$ 
10  $\pi_s \leftarrow \text{GenerateNIZKP}(\text{a majority of } bal'_s \text{ in } results \text{ are consistent});$ 
11  $\pi_r \leftarrow \text{GenerateNIZKP}(\text{a majority of } bal'_r \text{ in } results \text{ are consistent});$ 
// Phase 3: Ordering & Phase 4: Validation
12  $tx \leftarrow tx \cup \{bal'_s, bal'_r, \pi_s, \pi_r\};$ 
13  $\text{SendForOrderingAndValidation}(tx);$ 

```

---

and  $val_r$ , using  $c_s$ 's and  $c_r$ 's encryption keys, respectively. In the execution phase, executors use  $val_s$  and  $val_r$  to perform arithmetic evaluation. The reason for generating two different ciphertexts for  $val$  is that GAFE employs single-key FHE, which allows computation on two ciphertexts only if they are encrypted using the same encryption key. GAFE does not use multi-key FHE, which supports computation on ciphertexts encrypted with different encryption keys, because multi-key FHE is currently prohibitively expensive and impractical for real-world applications [31].

**Phase 1.2: Proving preconditions.** In certain applications, the client generates precondition NIZKPs to prove the satisfaction of preconditions. The correct transaction execution of these applications is conditional upon on these conditions. For instance, digital payment demands that the sender client's balance  $bal_s$  must be greater than or equal to the payment amount  $val$ . To prove this condition,  $c_s$  generates an NIZKP  $\pi_{bal \geq val}$  that takes  $c_s$ 's decryption key as private input, decrypts the two ciphertexts  $bal_s$  and  $val_s$ , and compares the resulting plaintexts. Thanks to the zero-knowledge properties of NIZKPs, GAFE ensures the preservation of privacy for  $c_s$ 's decryption key and the two plaintexts.

**Phase 1.3: Submitting transactions.** As the final step of the construction phase, the client submits the transaction, including the encrypted data and precondition NIZKPs, to all executors for execution.

**Phase 2: Execution.** As shown in Figure 1, the execution phase comprises four sub-phases. Algorithm 2 outlines Phase 2.1, 2.2, and 2.3 from the executor's viewpoint.

**Algorithm 2:** Execution phase of the executor.

---

```

1 buffer ← ∅;
// Phase 2.1: Verifying preconditions
2 Upon reception of transaction tx from client; do
3   {vals, valr, πbal≥val} ← tx;
4   bals ← GetBalance(ids); eks ← GetEvaluationKey(ids);
5   balr ← GetBalance(idr); ekr ← GetEvaluationKey(idr);
6   if VerifyNIZKP(πbal≥val, bals, vals) = true then
7     | buffer ← buffer ∪ {vals, valr, bals, balr, eks, ekr};
8   else
9     | Terminate();
// Phase 2.2: Executing transactions with GPU-accelerated FHE
10 Upon timeout or |buffer| ≥ bufferSize
11   MoveFromMainMemoryToGPUMemory(buffer);
12   For every transaction tx in buffer; do in parallel on GPU
13     | bal's ← FHESub(bals, vals, eks);
14     | bal'r ← FHEAdd(balr, valr, ekr);
15     | buffer ← buffer ∪ {tx, bal's, bal'r};
16   MoveFromGPUMemoryToMainMemory(buffer);
17   buffer ← ∅;
// Phase 2.3: Returning transaction results
18 For every transaction tx in buffer; do in parallel
19   | ReturnResultToClient(tx, ids)

```

---

**Phase 2.1: Verifying preconditions.** Upon receiving a transaction from a client, the executor initially verifies the validity of the precondition NIZKPs associated with the transaction, such as  $\pi_{bal \geq val}$ . GAFE proceeds to the subsequent phases iff  $\pi_{bal \geq val}$  is deemed valid. Otherwise, GAFE immediately terminates the transaction execution.

**Phase 2.2: Executing transactions with GPU-accelerated FHE.** Similar to Hyperledger Fabric, each GAFE executor independently executes transaction. For the on-GPU FHE computation, the executor performs concurrency control on the transaction level. Internally, the executor maintains a buffer that temporarily store the transactions received within a predefined time frame (e.g., 500 milliseconds). When a timeout occurs or the number of buffered transactions exceeds a specific threshold, the executor performs the following three steps: (1) moving the encrypted data of all buffered transactions from main memory to GPU memory, (2) launching the corresponding GPU kernel of FHE computation, and (3) copying back the resulting ciphertexts back to main memory. Taking the example of digital payment, the executor first moves the following data to GPU memory: the two ciphertexts representing the payment amount  $\{val_s, val_r\}$ , and the balances of the sender and the receiver clients  $\{bal_s, bal_r\}$ . Next, the executor

launches the GPU kernel for FHE addition to compute the updated balances:  $bal'_s = bal_s - val_s$  for the sender, and  $bal'_r = bal_r - val_r$  for the receiver. After completing the FHE computation, the executor copies back the updated balances  $\{bal'_s, bal'_r\}$  to main memory.

**Phase 2.3: Returning transaction results.** After the completion of transaction execution, the executor returns the transaction results to the client who initiated the transaction. In the case of digital payment transactions, the executor returns the updated balances  $bal'_s, bal'_r$  to the sender client  $c_s$ .

**Phase 2.4: Proving execution correctness.** Upon receiving the results of a transaction from all executors, the involving clients of the transaction are responsible for proving the execution correctness. To achieve this, the involving clients generates NIZKPs to prove the consistency of the majority of the results. In the case of digital payment, the sender  $c_s$  generates an NIZKP  $\pi_s$  that takes  $c_s$ 's decryption key as private input, decrypts  $c_s$ 's updated balance ciphertexts from all the results, and checks the consistency of the resulting plaintexts. Similarly, to ensure the consistency of the receiver client  $c_r$ 's updated balance,  $c_r$  follows the same procedure and generates an NIZKP  $\pi_r$  using  $c_r$ 's decryption key. Lastly,  $c_s$  sends the necessary materials to the orderers for ordering, including the transaction data  $\{val_s, val_r\}$ , the precondition NIZKP  $\pi_{bal \geq val}$ , the consistent results  $\{bal_s, bal_r\}$ , and the correctness NIZKPs  $\{\pi_s, \pi_r\}$ .

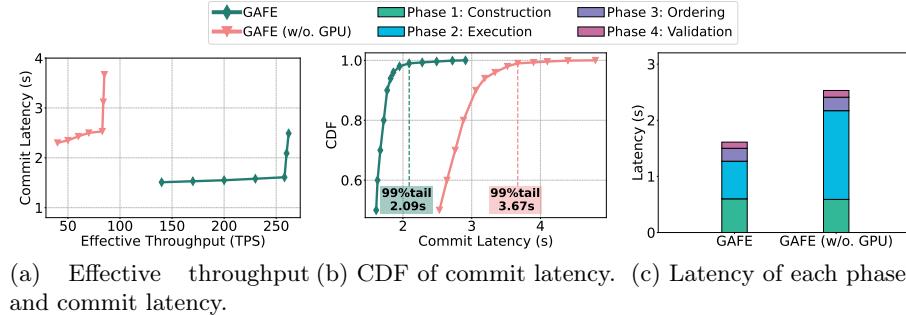
**Phase 3: Ordering.** GAFE orderers run an instance of a BFT consensus protocol, such as BFT-SMaRT [4]), to collectively determine the transaction order within each block. Once a consensus is reached among all orderers, they proceed to generate the block and disseminate the block to all executors for validation.

**Phase 4: Validation.** Upon receiving a block from the orderers, the executor follows the pre-determined order to sequentially validate each transaction within the block. The executor commits only those transactions that meet two conditions. First, the transaction must not have any write-conflicts with previously committed transactions within the same block. This prevents conflicts between concurrent transactions. Second, the transaction must be accompanied by valid correctness NIZKPs  $\{\pi_s, \pi_r\}$ . These NIZKPs serve as proofs of the transaction's execution correctness. If a transaction fails to meet either of these conditions, the executor aborts the transaction. Once the executor has validated all transactions, the executor permanently appends the block to the local copy of the blockchain.

## 5 Evaluation

### 5.1 Settings

**Implementation.** We built a prototype system of GAFE based on Hyperledger Fabric [2] and simulates the scenario of digital payment. We implemented the CKKS scheme [10] for GAFE based on the state-of-the-art studies on GPU-accelerated FHE schemes [22,28,30]. GAFE adopted the gnark [11] library's implementation for the Groth16 NIZKP system [13] and employed the BFT-SMaRT [4] consensus protocol. To evaluate the performance gain of GAFE, we



(a) Effective throughput and commit latency. (b) CDF of commit latency. (c) Latency of each phase.

Fig. 2: End-to-end performance of GAFE and the baseline.

also developed a baseline system called GAFE (w/o. GPU). The baseline follows a similar transaction execution workflow as GAFE, except that the baseline does not buffer transactions for concurrent execution and performs all FHE computations exclusively on the CPU.

**Metrics.** We evaluated two metrics: (1) *effective throughput*, which represents the average number of client transactions per seconds (TPS) committed to the blockchain; and (2) *commit latency*, which measures the time duration from transaction construction to commitment. Additionally, we also reported the cumulative distribution function (CDF) of commit latency for all committed transactions, and the latency of each phase in the transaction execution workflow (§4.2).

**Testbed.** We ran all evaluations on a cluster consisting of 4 machines. Each machine was equipped with an Nvidia RTX 3090 GPU, a 3.1GHz AMD EPYC 9754 CPU, and 64GB of main memory.

## 5.2 End-to-End Performance

We evaluated the end-to-end performance of GAFE and GAFE (w/o. GPU). For each evaluation, we initially instantiated three executors, four orderers, and one thousand clients. Next, we construct and submit 100,000 digital payment transactions using Tape [15], an efficient benchmark tool for blockchain. We explicitly ensured that no two transactions shared identical involving clients in order to prevent transaction aborts caused by write conflicts. We ran the evaluation ten times and reported the average values of the metrics.

GAFE exhibited exceptional end-to-end performance, as shown in Figure 2a. GAFE achieved a high throughput of 258 TPS and a low average latency of 1.61 seconds. In contrast, GAFE (w/o. GPU) displayed a significantly lower average throughput of 83 TPS and a notably longer average latency of 2.53 seconds. These results highlight the substantial performance advantage of GAFE over GAFE (w/o. GPU), with a 3.1× increase in effective throughput and a 37% reduction in commit latency. Additionally, Figure 2b illustrates that GAFE achieved a shorter 99% tail latency (2.09 seconds) compared to the baseline (3.67

seconds). This suggests that, even in worst-case scenarios, GAFE is capable of completing transaction execution in significantly less time.

GAFE’s high performance is attributed to the concurrent execution of FHE computation on GPU. Figure 2c compares the latency of each phase between GAFE and the baseline. Note that GAFE achieved significantly lower latency in the execution phase while achieving similar latencies in other phases. This latency reduction is made possible by the optimized parallel processing capability of the GPU, allowing for the concurrent execution of a substantial portion of the arithmetic computations involved in typical FHE schemes. As a result, GAFE avoids performing FHE computation on the CPU, which has fewer cores and is less efficient in executing a large number of compute-intensive computations in parallel. This capability enables GAFE to achieve higher effective throughput and shorter commit latency.

## 6 Conclusion

We present GAFE, a privacy-preserving blockchain that achieves high performance via the novel GPU-accelerated transaction execution workflow. GAFE protects data privacy through the use of FHE-encrypted transaction data, ensures the correctness of execution results by generating NIZKPs, and obtains high performance by leveraging GPUs to execute transactions concurrently. We implemented GAFE on top of Hyperledger Fabric and the latest advancements in GPU-accelerated FHE. Our evaluations demonstrate the superior performance of GAFE compared to the baseline, with a significant  $3.1\times$  increase in effective throughput (258 TPS) and a notable 37% decrease in commit latency (1.61 seconds).

## References

1. Acar, A., Aksu, H., Uluagac, A.S., Conti, M.: A survey on homomorphic encryption schemes: Theory and implementation. *ACM Computing Surveys (Csur)* **51**(4), 1–35 (2018)
2. Androulaki, E., Barger, A., Bortnikov, V., Cachin, C., Christidis, K., De Caro, A., Enyeart, D., Ferris, C., Laventman, G., Manevich, Y., et al.: Hyperledger fabric: a distributed operating system for permissioned blockchains. In: Proceedings of the thirteenth EuroSys conference. pp. 1–15. ACM, New York, NY, USA (2018)
3. Bessani, A., Sousa, J., Alchieri, E.E.: State machine replication for the masses with bft-smart. In: 2014 44th Annual IEEE/IFIP International Conference on Dependable Systems and Networks. pp. 355–362. IEEE, IEEE Computer Society, 1730 Massachusetts Ave., NW Washington, DCUnited States (2014)
4. Bessani, A., Sousa, J., Alchieri, E.E.: State machine replication for the masses with bft-smart. In: 2014 44th Annual IEEE/IFIP International Conference on Dependable Systems and Networks. pp. 355–362 (2014). <https://doi.org/10.1109/DSN.2014.43>
5. Brakerski, Z.: Fully homomorphic encryption without modulus switching from classical gapsvp. In: Advances in Cryptology–CRYPTO 2012: 32nd Annual Cryptology

- Conference, Santa Barbara, CA, USA, August 19-23, 2012. Proceedings. pp. 868–886. Springer, Springer Berlin, Heidelberg (2012)
6. Brakerski, Z., Gentry, C., Vaikuntanathan, V.: (leveled) fully homomorphic encryption without bootstrapping. ACM Transactions on Computation Theory (TOCT) **6**(3), 1–36 (2014)
  7. Bünz, B., Agrawal, S., Zamani, M., Boneh, D.: Zether: Towards privacy in a smart contract world. In: International Conference on Financial Cryptography and Data Security. pp. 423–443. Springer (2020)
  8. Castro, M., Liskov, B., et al.: Practical byzantine fault tolerance. In: OSDI. vol. 99, pp. 173–186. ACM, New York, NY, USA (1999)
  9. Cheng, R., Zhang, F., Kos, J., He, W., Hynes, N., Johnson, N., Juels, A., Miller, A., Song, D.: EKiden: A platform for confidentiality-preserving, trustworthy, and performant smart contracts. In: 2019 IEEE European Symposium on Security and Privacy (EuroS&P). pp. 185–200. IEEE (2019)
  10. Cheon, J.H., Kim, A., Kim, M., Song, Y.: Homomorphic encryption for arithmetic of approximate numbers. In: Advances in Cryptology–ASIACRYPT 2017: 23rd International Conference on the Theory and Applications of Cryptology and Information Security, Hong Kong, China, December 3–7, 2017, Proceedings, Part I 23. pp. 409–437. Springer (2017)
  11. consensys: Gnark (2023), <https://docs.gnark.consensys.net/>
  12. Fan, J., Vercauteren, F.: Somewhat practical fully homomorphic encryption. Cryptology ePrint Archive, Paper 2012/144 (2012), <https://eprint.iacr.org/2012/144>, <https://eprint.iacr.org/2012/144>
  13. Groth, J.: On the size of pairing-based non-interactive arguments. In: Fischlin, M., Coron, J.S. (eds.) Advances in Cryptology – EUROCRYPT 2016. pp. 305–326. Springer Berlin Heidelberg, Berlin, Heidelberg (2016)
  14. Hopwood, D., Bowe, S., Hornby, T., Wilcox, N., et al.: Zcash protocol specification. GitHub: San Francisco, CA, USA **4**(220), 32 (2016)
  15. Hyperledger-TWGC: Hyperledger-twgc/tape: A simple traffic generator for hyperledger fabric (2023), <https://github.com/Hyperledger-TWGC/tape>
  16. Kalodner, H., Goldfeder, S., Chen, X., Weinberg, S.M., Felten, E.W.: Arbitrum: Scalable, private smart contracts. In: 27th USENIX Security Symposium (USENIX Security 18). pp. 1353–1370 (2018)
  17. Kang, H., Dai, T., Jean-Louis, N., Tao, S., Gu, X.: Fabzk: Supporting privacy-preserving, auditable smart contracts in hyperledger fabric. In: 2019 49th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN). pp. 543–555. IEEE, IEEE, Portland, Oregon, USA (2019)
  18. Kosba, A., Miller, A., Shi, E., Wen, Z., Papamanthou, C.: Hawk: The blockchain model of cryptography and privacy-preserving smart contracts. In: 2016 IEEE symposium on security and privacy (SP). pp. 839–858. IEEE (2016)
  19. Meier, A.V.: The elgamal cryptosystem. In: Joint Advanced Students Seminar (2005)
  20. Monero: The monero project, <https://www.getmonero.org/>
  21. Nakamoto, S.: Bitcoin: A peer-to-peer electronic cash system (2008)
  22. Özcan, A.Ş., Ayduman, C., Türkoğlu, E.R., Savaş, E.: Homomorphic encryption on gpu. IEEE Access (2023)
  23. Patel, R., Migliavacca, M., Oriani, M.E.: Blockchain in banking and finance: A bibliometric review. Research in International Business and Finance **62**, 101718 (2022)

24. Raczyński, M.: What is the fastest blockchain and why? analysis of 43 blockchains (Jan 2021), <https://alephzero.org/blog/what-is-the-fastest-blockchain-and-why-analysis-of-43-blockchains/>
25. Steffen, S., Bichsel, B., Baumgartner, R., Vechev, M.: Zeestar: Private smart contracts by homomorphic encryption and zero-knowledge proofs. In: 2022 IEEE Symposium on Security and Privacy (SP). pp. 179–197. IEEE (2022)
26. Sun, X., Yu, F.R., Zhang, P., Sun, Z., Xie, W., Peng, X.: A survey on zero-knowledge proof in blockchain. IEEE network **35**(4), 198–205 (2021)
27. Viand, A., Jattke, P., Hithnawi, A.: Sok: Fully homomorphic encryption compilers. In: 2021 IEEE Symposium on Security and Privacy (SP). pp. 1092–1108. IEEE, San Francisco, CA, USA (2021). <https://doi.org/10.1109/SP40001.2021.00068>
28. Wang, Z., Li, P., Hou, R., Li, Z., Cao, J., Wang, X., Meng, D.: He-booster: An efficient polynomial arithmetic acceleration on gpus for fully homomorphic encryption. IEEE Transactions on Parallel and Distributed Systems **34**(4), 1067–1081 (2023)
29. Xu, L., Zhang, Y., Zhu, L.: Regulation-friendly privacy-preserving blockchain based on zk-snark. In: International Conference on Advanced Information Systems Engineering. pp. 167–177. Springer (2023)
30. Yang, H., Shen, S., Dai, W., Zhou, L., Liu, Z., Zhao, Y.: Implementing and benchmarking word-wise homomorphic encryption schemes on gpu. Cryptology ePrint Archive (2023)
31. Yuan, M., Wang, D., Zhang, F., Wang, S., Ji, S., Ren, Y.: An examination of multi-key fully homomorphic encryption and its applications. Mathematics **10**(24), 4678 (2022)