HEOV: Achieving High-Performance Confidential Permissioned Blockchain via Correlated-Merged Homomorphic Encryption

Abstract

Permissioned blockchains enable a wide range of applications that emphasize data confidentiality (e.g., healthcare). Unfortunately, existing confidential permissioned blockchains do not efficiently utilize compute-intensive cryptographic primitives, leading to low performance.

We present HEOV, a high-performance confidential permissioned blockchain. We incorporate fully homomorphic encryption (FHE) to protect data confidentiality and enable direct arithmetic computation over ciphertext, eliminating the overhead of data decryption and re-encryption required by smart contract execution. To reduce the invocation of costly cryptographic operations, we propose a novel correlated-merging protocol that detects and merges multiple mergeable transactions as one. To verify result correctness, we develop a non-interactive zero-knowledge proof (NIZKP) for each transaction. By leveraging permissioned blockchains' endorsement mechanism, we tailor NIZKPs to prove the equality of peers' individual execution result, rather than re-executing smart contracts. This approach greatly simplifies the generation and verification of NIZKPs. Extensive evaluation results show that without loss of data confidentiality, HEOV gains exceptional performance with an increase of up to $7.14\times$ in effective throughput and a decrease of 13% in average end-to-end latency, rendering it suitable for deploying diverse blockchain applications.

1 Introduction

A permissioned blockchain is an immutable ledger among explicitly identified mutually untrusted participant clients and nodes. Unlike permissionless blockchains, the permissioned blockchain is decoupled from the cryptocurrency [35] and can run traditional high-performance Byzantine-fault-tolerant (BFT) consensus protocols (e.g., BFT-SMaRT [7]) to tolerate malicious participants. Therefore, permissioned blockchains are ideal for deploying performance-sensitive and security-critical enterprise applications, such as citizen information


Figure 1: The EOV workflow

management [31] and medical chain [32].

The execute-order-validate (EOV) workflow, introduced by Hyperledger Fabric (HLF) [3], is one of the most notable permissioned blockchain workflows [3, 25, 40]. As illustrated in Figure 1, the EOV workflow processes client transactions in three phases. First, in the *execute* phase, the client submits transactions to the executor nodes for policy-based executions. A transaction's execution result is correct if certain executor nodes (specified in the application-layer endorsement policy [3]) execute the transaction and produce consistent execution results. Then, in the *order* phase, orderer nodes reach a consensus on the order of execution results. Finally, in the *validate* phase, executor nodes validate the ordered transactions and their results, commit valid ones, and abort the invalid ones.

Two advancements in the EOV workflow make it especially suitable for deploying enterprise applications. First, the EOV workflow concurrently executes all transactions to achieve a high performance [3, 40, 43]. Second, the consensus-on-result feature enables EOV to tolerate non-deterministic transactions, because all executors can commit a consistent execution result. Therefore, EOV can support applications written in general-purpose multi-threaded programming languages (e.g., Golang [24]). These unique advancements make EOV extensively used in various enterprise applications. For example, Singapore deployed a stock exchange system on HLF [3]; Walmart built its supply chain system [18] on HLF.

However, existing EOV permissioned blockchains often involve a trade-off in data privacy. Specifically, a client must submit each transaction's plaintext to the executor nodes for execution, which may jeopardize user data privacy. This is particularly problematic for applications that handle sensitive data, such as patients' medical records on a medical chain [32]. Systems that handle such sensitive data are typically subject to strict privacy laws and regulations, such as the General Data Protection Regulation in Europe [52]. Therefore, ensuring data privacy in EOV permissioned blockchains is a critical issue that needs to be resolved.

Recent studies [30, 44, 46–48] have revealed the potential of homomorphic encryption (HE) and non-interactive zero-knowledge proof (NIZKP) as powerful tools for safeguarding data privacy in blockchain systems. The notable study ZeeStar [46] leverages additive HE and NIZKP to enable encrypted, privacy-preserving transaction executions among mutually untrusted participants. Specifically, blockchain nodes first execute transactions on data owned by different participants using HE, and then generate NIZKP to verify the correctness of the executions on other blockchain nodes.

Unfortunately, this innovative approach still faces the dilemma between the generality of supported applications and achieving high performance. On the one hand, adopting the fully homomorphic encryption (FHE [51]), which offers both addition and multiplication operations, can support general blockchain applications. However, to prove the correctness of transaction executions, FHE leads to intricate NIZKP and prohibitively high costs for generating these proofs. For instance, generating a Groth16 zk-SNARK for Paillier encryption requires over 256 GB of RAM, rendering it impractical for commodity desktop machines [46]. Moreover, generating a single proof takes several seconds. Such prolonged latency is unacceptable for blockchain applications, which generally require a commit latency of a few seconds [46]. On the other hand, partial homomorphic encryption (PHE) schemes offer significant advantages over FHE schemes in terms of performance. According to a recent survey [15], a simple proof-of-concept implementation of the ElGamal scheme [33], operating at a 128-bit security level, requires less than 20 microseconds to perform a homomorphic addition. In contrast, a highly optimized FHE library such as Microsoft SEAL [42] takes at least 500 milliseconds to finish the same arithmetic operation, resulting in a $25\times$ performance gap. However, PHE hampers the support for multiplications of data owned by different participants, thus limiting the system's generality. Overall, existing approaches can achieve only either high performance or application generality but not both.

In this study, our key insight to tackle the dilemma is that, we can leverage EOV's policy-based execution, which enables the blockchain nodes to efficiently verify the correctness of transaction executions, to obviate the necessity of verifying the correctness of executions using the NIZKP. With EOV's policy-based execution, each transaction is homomorphically executed on multiple executor nodes, and the correctness of the executions is validated by comparing the consistency of execution results from executor nodes. As a result, the system can seamlessly integrate with FHE schemes to support a wide array of general blockchain applications while still maintaining high performance.

This insight leads to HEOV, the first high-performance privacy-preserving EOV permissioned blockchain framework that supports general blockchain applications. HEOV clients leverage FHE to encrypt sensitive user data, while executor nodes execute transactions on encrypted ciphertext. By synergizing the strengths of EOV and HE, HEOV ensures the correctness of transactions' executions by embracing EOV's policy-based executions, and preserves data confidentiality by preventing executors from accessing sensitive user data in plaintext.

The integration of the EOV workflow with FHE presents two unique challenges. The first challenge is tolerating non-deterministic transactions in EOV permissioned blockchains. For a non-deterministic transaction, different executors may produce inconsistent execution results, and EOV only commits transactions with consistent execution results. However, in HEOV, the executors cannot directly compare the FHE-encrypted execution results like existing EOV permissioned blockchains, due to the random noise that FHE inserts into the ciphertexts to defend against ciphertext attacks [21]. Consequently, the executors cannot determine whether a transaction is non-deterministic or not by checking the consistency of the transaction's execution results, as the ciphertexts of execution results are inconsistent even for a deterministic transaction.

The second challenge is the potentially disastrous performance degradation when integrating FHE with EOV, compared to plain text executions. First, FHE schemes incur a high computation overhead. Existing studies [51] have reported that FHE-encrypted transactions take 10^6x longer time to execute, which we confirmed in our evaluation (§7). Worse, the random noise that FHE inserts into ciphertext accumulates as the number of homomorphic computations increases. The ciphertext can be only decrypted when the noise is below a threshold, limiting the number of consecutive HE computations (e.g., multiplication) on the same ciphertext before decryption fails. This issue can be addressed by resetting the noise using relinearization, which can take tens of seconds even in efficient implementations [51].

This performance issue is further exacerbated by EOV's conflicting aborts. Due to EOV's concurrent transaction execution, when multiple concurrently executed transactions access the same key and at least one of them writes to the key (conflicting transactions), only one transaction can commit and other transactions will abort. Such conflicting transactions are prevalent in typical blockchain applications [43]. For example, HLF achieved only 24.5% of its conflict-free peak throughput in the stock trading application, as confirmed in Figure 4. These two factors are intertwined, as conflicting transactions

result in a large number of HE-encrypted ciphertext computations on the same key, leading to frequent relinearizations.

To tackle the first challenge, HEOV clients generate NIZKP to prove the consistency of execution results of different executor nodes. Instead of generating heavyweight NIZKP to prove the correctness of the entire transaction execution process as existing studies [46], HEOV generates only lightweight NIZKP to prove the results' plaintext consistency without introducing prohibitive proof generation overhead. This enables HEOV to integrate FHE for supporting general blockchain applications with both addition and multiplication operations. Furthermore, the NIZKP safeguards against malicious participants falsifying the consistency of results, ensuring compatibility with BFT safety.

To tackle the second challenge, our basic idea is to exploit the conflicting transactions, which are ubiquitous in permissioned blockchain applications, to reduce the invocation number of homomorphic computations and relinearizations. We propose a new secure computing abstraction called correlated-merged homomorphic encryption for HEOV. This approach enhances the online schedule of transactions with an offline analysis of the transaction code. During the offline analysis, HEOV analyzes the smart contract and identifies whether multiple conflicting transactions invoking it can be merged into a single transaction. Based on the results of the offline analysis, HEOV's online schedule merges conflicting transactions and performs the necessary HE-encrypted ciphertext computations on the merged transactions. For example, when two transactions both add to the same key (e.g., two $X + 1$ operations), HEOV combines the two adding transactions into a single add transaction on that key (e.g., one $X + 2$ operation). This reduces the number of HE computations and relinearization required while preserving the semantics of the original transactions. Moreover, in contrast to existing permissioned blockchains with EOV consensus (e.g., HLF [3]), where conflicting transactions are resolved by committing only one and invalidating the rest, HEOV can commit all conflicting transactions without any aborts.

We implemented HEOV on the codebase of HLF, and compared HEOV with state-of-the-art privacy-preserving blockchain systems [46]. Our evaluation and analysis results show that:

- HEOV is efficient. Its throughput was up to 7.14x higher than our baselines (§7).

- HEOV achieved even higher throughput as the conflict ratio increased (§7.1).

- HEOV is robust to no-resubmission attack (§4).

HEOV makes a pioneering effort in identifying the challenges posed by the integration of HE and EOV, achieving the best of both worlds. HEOV integrates EOV with HE to achieve data confidentiality without introducing significant performance overhead or sacrificing the generality in supporting diverse blockchain applications. The implication behind HEOV 's synergy of conflicting transactions and HE extends beyond the scope of the blockchain and can serve as a general transaction scheduling mechanism for other exciting research areas for distributed systems, such as the Internet of Things [54] and databases [50], where transactions often conflict, and each execution of the distributed transaction commitment on the Internet leads to substantial latency.

2 Background And Related Works

In this section, we present a background overview of HEOV and review related works in the field of privacy-preserving blockchains, which are compared in Table 1.

2.1 Execute-Order-Validate Permissioned Blockchain

Permissioned [3, 14] and permissionless [23, 45] blockchains are maintained by multiple mutually untrusted organizations, but they differ significantly in their protocols. In permissioned blockchains, nodes are incentivized with cryptocurrency to follow the protocol, which requires specific proofs. For example, proof-of-work [35] involves high computation power, while proof-of-stake [22, 53] requires a large cryptocurrency holding. In contrast, permissioned blockchains employ a BFT consensus protocol among consensus nodes to tolerate malicious nodes instead of proof-of-work or stake. Permissioned blockchains, such as HLF and Cosmos, demonstrate better performance due to BFT protocols like PBFT [11] and Hot-Stuff [1]. HEOV, inherited from HLF, can concurrently run multiple orderer nodes forming a BFT ordering service that can determine the order of transactions for each block within a few hundred milliseconds.

Permissioned blockchains can be further categorized into two sub-types based on their workflows: order→execute (OE) and execute→order→validate (EOV), as discussed in §1. HEOV adopts the EOV workflow for parallelism and applies the correlated-merging protocol (i.e., pre-submission transaction merging) to minimize potential conflicting aborts, laying a solid foundation for the high performance of HEOV.

2.2 Homomorphic Encryption

Homomorphic encryption (HE) enables computation over encrypted data without revealing the plaintext. It is based on the concept of homomorphism, which is a structure-preserving map between algebraic structures. For instance, if $f : X \to Y$ is a map between two sets X and Y equipped with the same structure, and $\oplus$ is an operation of the structure (suppose $\oplus$ is a binary operation for simplicity), then f satisfies the property

$$f(a \oplus b) = f(a) \oplus f(b)$$

System	Data Privacy	Computation On Ciphertext	Computation On Foreign Data	Tolerate Non-Deterministic Smart Contract
HLF [3]	✗	✗	✓	✓
ZeeStar [46]	✓	✓	✓	✗
Zapper [48]	✓	✗	✗	✗
smartFHE [44]	✓	✓	✗	✗
FabZK [30]	✓	✗	✓	✗
HEOV	✓	✓	✓	✓

Table 1: Compare HEOV with existing privacy-preserving blockchain systems.

for every pair of elements a, b of X, and $a \oplus b$ of Y.

There are two major types of HE schemes: partially homomorphic encryption (PHE) and fully homomorphic encryption (FHE). PHE supports a specific operation (e.g., addition or multiplication), while FHE allows arbitrary combinations of computations. Despite simplicity, PHE is restricted by its limited computational capability and thus is not adopted by HEOV. For example, the Paillier cryptosystem [38, 39] only supports homomorphic addition. In contrast, FHE schemes (e.g., BFV [9, 19], BGV [10], and CKKS [12]) strike a balance between security and performance. HEOV employs the lattigo [34] library's BFV implementation, which can perform arithmetic computations on 58-bit unsigned integers within tens of milliseconds, achieving a good trade-off between security and performance.

2.3 Non-Interactive Zero-Knowledge Proofs

Non-interactive zero-knowledge proof (NIZKP) [20, 26] is a cryptographic primitive that enables a prover P to prove knowledge of a secret s to a verifier V without revealing s. Once P generates an NIZKP using a proving key, V can verify the proof using the corresponding verifying key without P being present. Formally, given a proof circuit ϕ, a private input s, and some public input x, prover P can use an NIZKP to prove knowledge of s satisfying a predicate $\phi(s;x)$. One popular type of NIZKP is called zero-knowledge succinct non-interactive arguments of knowledge (zkSNARKs) [8, 27, 55], which allows any arithmetic circuit ϕ and guarantees constant-cost proof verification in the size of ϕ, making them suitable for blockchain applications [2, 4, 16, 37, 47].

HEOV employs the gnark library [13]'s Groth16 implementation, a zkSNARK construction known for its constant size of proof (several kilobytes) and short verifying time (tens of milliseconds), for two important reasons. Firstly, NIZKP is necessary for enabling conditional branching in the context of encrypted data. Since direct comparison between two ciphertexts is currently infeasible, traditional if-statements commonly used in programming languages is not applicable. NIZKP provides a solution to this problem. Secondly, NIZKP is essential to prove the correctness of transaction results generated by various executors, thereby preventing malicious clients from forging transaction results.

2.4 Related Works

Smart Contract Privacy. Privacy is crucial in smart contracts, and previous works have achieved varying levels of privacy.

SmartFHE [44] and ZeeStar [46] focus on *data privacy*, where only those with the decryption key, typically the data owner, can access the specific data. SmartFHE proposes a framework using FHE and NIZKP, while ZeeStar adopts El-Gamal [33] as an additive homomorphic scheme and develops a Solidity-based domain-specific language and transpiler for private smart contracts.

Zapper [48] and FabZK [30] go beyond *data privacy* and also hide the data's identity, known as *key privacy*. Zapper achieves *key privacy* through a combination of an oblivious Merkle tree construction and an NIZK processor, but it sacrifices smart contract expressiveness. FabZK achieves *key privacy* by using zero-knowledge proof, verifiable Pedersen commitments, and a shared tabular structured ledger that stores only encrypted data and conceals the relevant members of transactions.

Smart Contract Generality. Previous works have achieved different levels of smart contract generality using different approaches, as compared in Table 1.

ZeeStar supports general-purpose Ethereum [53] smart contracts but has limited functionality. It only allows additive homomorphic encryption and partially implements multiplication for few limited scenarios. Additionally, it uses NIZKP to prove the correctness of offline HE computation, which requires significant proof generation time, as confirmed in its evaluation. These seriously undermine the applicability of ZeeStar. In contrast, HEOV addresses these two pain points by adopting FHE and light-weight NIZKP.

Zapper sacrifices smart contract generality by prohibiting control flow and computation on data belonging to multiple users. SmartFHE currently implements a single-key variant where all private inputs for a transaction must be owned by the same party, while the multi-key variant is currently impractical, leading to the same problem that hinders Zapper's generality. FabZK achieves stronger generality but faces scalability issues because every transaction causes an update of the values of all organizations, which puts heavy pressure on performance and even complicates the process of managing the number of FabZK organizations on-the-fly.

3 Overview

3.1 System Model

HEOV consists of three types of participants: clients, executors, and orderers. The latter two are often referred to as nodes because they act as "servers" in an HEOV network with which clients interact. An HEOV network is typically maintained by multiple organizations, each consisting of clients and nodes. Members of the same organization are mutually trusted, but those from various organizations do not trust each other.

Clients. Clients are end-users who submit transaction proposals to executors for execution and orderers for ordering. Additionally, clients are responsible for generating NIZKP to prove the consistency of transaction results.

Executors. All executors in HEOV are responsible for evaluating transactions and maintaining a local copy of the blockchain, similar to HLF. HEOV executors utilize libHEOV (§6.1), which offers a set of APIs that support NIZKP and HE operations to perform the execution of transactions with homomorphically encrypted data. libHEOV is tailored for EOV contract execution and comes with pre-defined parameters for HE; therefore, smart contract developers do not need to be cryptography experts to efficiently incorporate smart contracts with HE operations.

Besides, a *lead executor* is elected for each organization to implement the correlated-merging protocol, as shown in Figure 2. Lead executors are not different from other executors in transaction execution and blockchain maintenance, but they are assigned the following extra tasks.

- Receive transaction proposals from clients of its organization.

- Optimize transactions by merging multiple transactions targeting the same set of keys into a single transaction.

- Dispatch transaction proposals to other executors and orderers for execution and ordering, respectively.

- Pass transaction responses to clients.

Clients are not required to directly submit their transaction proposals to a lead executor as other executors will finally forward the proposals to it.

Orderers. Orderers in HEOV are the same as in HLF. Specifically, they determine the order of transactions via some consensus protocol (e.g., BFT) and pack them into individual blocks, which are ultimately validated and recorded by relevant executors.

HEOV identifies all participants through a membership mechanism similar to HLF, where each participant is issued a pair of public/secret keys that encode their identity information. This mechanism enables access control to HEOV resources and, more importantly, makes all operations on HEOV auditable and traceable since any transaction is signed by a specific participant with the corresponding privilege.

3.2 Threat model

HEOV adopts the Byzantine failure model [6, 11], where malicious orderers run BFT consensus protocols that tolerate up to f malicious orderers out of $3f + 1$ total orderers.

HEOV groups participants (i.e., clients, executor nodes, and orderer nodes) into organizations. Participants are mutually trusted only when they belong to the same organization. This implies that clients and nodes are mutually untrusted, as an organization only contains clients or nodes. Any clients or nodes can be malicious, in which case they are called attackers.

3.3 Workflow Overview

The runtime workflow of HEOV can be summarized into several steps, as shown in Figure 2. A more comprehensive discussion of the runtime workflow can be found in §4.2.

Step 1: Correlated merging. The lead executor collects transaction proposals, merges consecutive *mergeable* proposals, and leaves non-mergeable proposals as is.

Step 2: encryption. The lead executor encrypts the merged proposals using relevant organizations' public keys and signs each transaction with its certificate.

Step 3: dispatch. The lead executor dispatches the signed transactions to other executors, gathering clients' signatures and notifying clients of the transaction ID.

Step 4: execution. Executors invoke smart contracts, sign proposals, and send them back to the lead executor along with the result.

Step 5: proof generation. The lead executor generates an NIZKP to prove the plaintext consistency of results without re-running the smart contract.

Step 6: consensus. HEOV orderer nodes use a BFT protocol to reach a consensus on the order of transactions in a block.

Step 7: commit. Executors commit a transaction after receiving a valid block, ensuring no overlap in read-write set with previous valid transactions.

Overall, the highlight of the HEOV runtime workflow stems from achieving data privacy and high performance, which are ensured by several key elements. To guarantee the correctness of results, a lightweight NIZKP is generated. The workflow leverages the correlated-merging protocol, which optimizes mergeable transactions, effectively reducing the occurrence of conflicting transactions and the need for HE operations. To maintain data confidentiality, lead executors are responsible for safeguarding the private keys of their respective organizations. Additionally, the libHEOV library plays a crucial role in facilitating HE computations and NIZKP operations. By combining these elements, the HEOV runtime workflow achieves both data privacy and high performance for permissioned blockchain smart contracts.
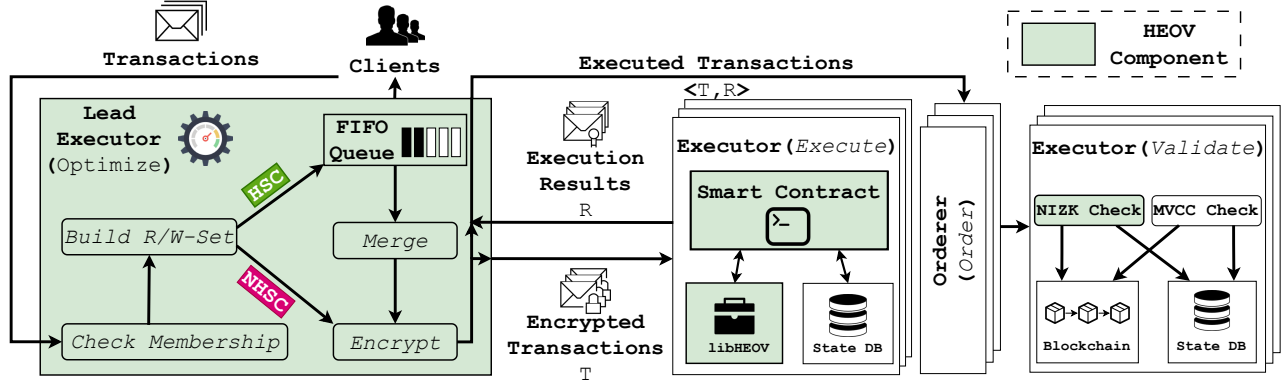
Figure 2: The HEOV Runtime Workflow. HEOV components are highlighted in green. An "HSC" transaction invokes a Homomorphic Smart Contract, while the "NHSC" one does not.

4 Protocol

## 4.1 Offline Homomorphic Smart Contract Analysis

Figure 3: Offline Analysis Protocol. "AH", "MH", and "NH" means additive, multiplicative, and non-homomorphic, respectively.

HEOV's offline analysis protocol determines if a smart contract is *homomorphic*, meaning multiple transactions modifying the same set of states can be merged without affecting the execution results. As shown in Figure 3, before deployment, a smart contract analyzer conducts the analysis and generates metadata indicating if the contract is additive homomorphic (**AH**), multiplicative homomorphic (**MH**), or non-homomorphic (**NH**). This metadata is essential for permitting lead executors to perform transaction merging.

HEOV analyzes smart contracts with random interpretation [2], which generates different input and initial states to capture all possible control flow paths. As a first step, two random input sets I_1 and I_2, along with a random initial memory state M, are generated for the candidate smart contract SC. The memory state contains the SC outputs and all keys written by SC. Then, SC is interpreted in two different ways. First, HEOV executes SC twice consecutively: first with I_1, producing M_{I_1}, and then with I_2, producing the final state M_{I_1,I_2}. Simultaneously, HEOV begins with the same initial state M and executes SC once with $I_1 + I_2$, producing $M_{I_1+I_2}$. If $M_{I_1,I_2} =$

$M_{I_1+I_2}$ for all control flow paths, SC is additive homomorphic. Similarly, if $M_{I_1,I_2} = M_{I_1 \times I_2}$, it is multiplicative homomorphic.

4.2 Runtime Transaction Execution and Commitment

The runtime transaction execution and commitment protocol specifies the rules for HEOV's transaction execution, as visualized in Figure 2.

Phase 1: Submission. In the runtime protocol, the client first submits a transaction proposal to an executor within its organization. If the receiving leader is not a lead executor, it forwards the proposal to the designated lead executor. The submission process uses mutual TLS, which ensures the confidentiality of proposal arguments and requires both parties to authenticate themselves by presenting certificates to establish their identities.

Phase 2: Pre-execution preparation. Upon receiving a transaction from a client, the lead executor undertakes a series of sequential preparation tasks: (**Phase 2.1**) constructing the read-write set, (**Phase 2.2 and 2.3**) merging homomorphic transactions, and (**Phase 2.4**) encrypting transaction arguments.

Phase 2.1: Read-write set construction. The read-write set of a transaction includes the read set and the write set, which encompass the states being read and written, respectively. Constructing a read-write set helps identify the owner of the value being written, which is necessary for selecting the appropriate public key to encrypt the value into ciphertext.

Phase 2.2: Transaction diversion. During Phase 2.2 and 2.3, lead executors handle transactions differently based on the homomorphic property of the smart contracts they invoke. This property is determined through the offline homomorphic smart contract analysis protocol mentioned in §4.1. For non-homomorphic smart contracts, the correlated-merging protocol is not applicable, and proposals for these contracts undergo direct encryption as specified in Phase 2.4. For homomorphic smart contracts, proposals are temporarily stored in

a FIFO queue, which collects all the proposals that will be packaged into the same block, thereby facilitating the potential merging of these proposals.

Phase 2.3: Correlated merging. The lead executor merges consecutive correlated proposals from the FIFO queue, specifically those that invoke the same homomorphic smart contract and have overlapping read-write sets. This merging process reduces the number of transaction proposals by merging certain proposals into newly merged ones. For example, consider two correlated *transfer* transactions, T_1 and T_2, with shared relevant parties and inputs I_1 and I_2, respectively. After merging, a new transaction T_{1+2} is created, incorporating the combined input $I_1 + I_2$, effectively replacing T_1 and T_2.

Phase 2.4: Homomorphic encryption. Before submission, the arguments of each proposal undergo encryption using public keys corresponding to the identities listed in the write set. This encryption step is crucial because each argument may be utilized on different ciphertexts owned by clients from other organizations. For example, if *Sam* transfers M tokens to *Rachel*, the plaintext value M must be encrypted into two ciphertext, $Ct_{M,1}$ and $Ct_{M,2}$, using the public keys of *Sam*'s and *Rachel*'s organizations, respectively.

Phase 3: Dispatch. The lead executor signs and dispatches the proposal to other executors based on the endorsement policy. The client transaction ID and proposal signatures are retained as metadata, enabling clients to verify the correct execution of their proposals.

Phase 4: Execution. An HEOV executor executes a smart contract and sends the execution response to the lead executor. Unlike HLF, HEOV smart contracts leverage the capabilities of libHEOV for HE and NIZKP operations. Additionally, HEOV smart contracts use NIZKP as an alternative to the *if* statements found in most programming languages. In HEOV, *if* statements are neither recommended nor allowed when ciphertext is involved in the condition expression because HE ciphertext does not support comparison operations such as $<$, $>$, or $=$, and executors are not authorized to access private keys, making it impossible to decrypt the ciphertext and compare the underlying plaintext. To address this, HEOV smart contracts generate NIZKP that validate the execution results and are stored on the blockchain for future reference.

Phase 5: Correctness NIZKP generation. Before submitting a transaction for ordering, the lead executor collects results from the executing executors. Each result comprises ciphertext and endorsement information (i.e., the acknowledgment of the transaction). The lead executor decrypts and compares the ciphertexts to ensure *plaintext consistency*. If a ciphertext belongs to another organization, the lead executor requests verification from that organization. Valid results require a quorum, typically a majority, of ciphertexts referencing the same plaintext. Only transactions with valid results are eligible for ordering. To ensure the integrity and prevent forgery of results, the lead executor generates an NIZKP to prove the plaintext consistency, which is also stored on the

blockchain and is publicly verifiable. By employing this approach, the lead executor cannot create a fake result or tamper with the on-chain data, as the verification process ensures the accuracy and integrity of the results.

Phase 6: Ordering. HEOV orderers run a BFT consensus protocol (e.g., BFT-SMaRT [7]) to reach a consensus on the order of transactions in one block. The block is then distributed to all executors for validation once finalized by orderers.

Phase 7: Validation. Upon receiving a block from orderers, an executor sequentially validates each transaction, including verifying the consistency NIZKP associated with the transactions. If validation is successful, the executor applies the transaction result to its local world state database, ensuring accurate reflection of transaction changes. After validating and processing all transactions in the block, the block is permanently appended to the local blockchain. Since all executors take deterministic and consistent actions for the same block, all local copies of the blockchain maintain consistency, ensuring the overall consistency of HEOV.

Phase 8: Response. As a final step, the lead executor responds to the client who proposed the transaction. This response informs the client whether their transaction has been accepted or rejected.

4.3 Defense Against No-Resubmission Attack

No-resubmission attack is a performance degradation attack that exploits the EOV workflow by maliciously skipping the transaction resubmission process. The attack is simple but detrimental because it wastes a significant amount of the computation power of executors, especially when smart contracts involve heavy HE computations. However, distinguishing no-resubmission attacks from legitimate actions can be challenging. A benign transaction can produce different results when executed by different executors, particularly for non-deterministic smart contracts. In such cases, the results of the current execution are discarded, and no-resubmission becomes a necessary action.

To mitigate the damages caused by no-resubmission attacks, HEOV adopts a per-executor counter-based approach. To be specific, each executor maintains a counter for each lead executor in the HEOV network, tracking the number of no-resubmission instances within a timeframe. If the count exceeds a configurable upper bound, the executor rejects transactions from that lead executor for a specified duration, mitigating the damages caused by no-resubmission attacks.

5 Analysis

This section defines and analyzes the three guarantees of HEOV: correctness, liveness, and confidentiality.

5.1 Correctness

Definition 1 (Correctness). *For any two benign nodes $Node^1$ and $Node^2$, they hold the same world state with the n-th block as the latest block. The $(n+1)$-th blocks received by $Node^1$ and $Node^2$, denoted as $Block^1_{n+1}$ and $Block^2_{n+1}$ respectively, are consistent and deterministic.*

The correctness of smart contract execution in HEOV is ensured by satisfying ACID and BFT safety properties at the transaction and block levels, respectively.

HEOV guarantees ACID properties for transaction execution, which are the foundation of all reliable failure-tolerant transaction processing systems and are indispensable as concurrent transaction execution is possible and common in EOV workflow.

- *Atomicity*. HEOV executes and commits (or aborts) each transaction as a whole.

- *Consistency*. HEOV only commits valid transactions with correct execution results (conform to the smart contract's logic) that satisfy the application-level endorsement policy 4.2. Specifically, the execution results of a transaction are considered valid as long as enough executors digitally sign the result.

- *Isolation*. HEOV inherits the isolation property from the EOV workflow. In other words, each HEOV transaction operates as if it is the only transaction executing, without interference or conflicts with other transactions.

- *Durability*. Same as existing EOV permissioned blockchains [3], HEOV inherits the durability guarantee from EOV's consensus protocol. Executors can retrieve the blockchain's transactions from their local copies of the blockchain.

HEOV achieves equivalent BFT safety as existing EOV blockchains [3] through the presence of two crucial properties:

- *Block consistency*. HEOV treats the consensus protocol as a blackbox, thereby enjoying the mature consistency (safety) guarantee of existing BFT consensus protocols [7] and their implementations. This ensures that all executors process and append the same set of blocks in an identical order determined by orderers.

- *State consistency*. Every executor follows the same deterministic protocol (§4.2) to validate each transaction in a given block. A transaction's result is committed to the state database *iff* all of its results produced by various executors are consistent and do not modify any key that has already been modified by a previous transaction in the same block. In this way, all executors maintain a consistent state and a local copy of the blockchain after validation of a newly appended block.

5.2 Liveness

Definition 2 (Liveness). *Let C be a benign client and T be a well-formed transaction submitted by C and accepted by the lead executor LE. T will eventually be committed to HEOV, even in the event of a crash of C or a reboot of LE.*

Two liveness requirements must be met to commit T to HEOV successfully.

Firstly, the orderers must include T in a publicly agreed-upon clock to guarantee its proper ordering and inclusion in the blockchain. This requirement is satisfied as HEOV inherits the BFT liveness guarantee from its BFT protocol [7] running in EOV's ordering phase.

Secondly, T needs to be submitted to and executed by executors during the execution phase. This is addressed by the combined liveness guarantees of HEOV lead executor (*LE*) and the inherited EOV workflow. After T is accepted by *LE*, it is stored in a persistent FIFO queue that withstands crashes. In the event of a crash, T is not lost and will be finally dispatched to executors after *LE* reboots. The endorsement policy [3], which is an application-level agreement protocol among participating organizations in the EOV workflow, ensures that a sufficient number of benign organizations execute the transactions. This prevents a certain number of malicious executors from tampering with the correct result. For example, in a typical token transfer application with three organizations (*orgA*, *orgB*, *orgC*), the endorsement policy is *majority(orgA, orgB, orgC)*, which requires that every T must be endorsed by at least two of them and ensures that at most one organization can act maliciously.

5.3 Confidentiality

Definition 3 (Confidentiality). *Suppose C is a smart contract following the guideline of §6.2. With this setup, data owners can securely store and process the ciphertext CT without compromising the confidentiality of the corresponding plaintext PT. An active attacker A cannot deduce PT from CT (storage confidentiality). Moreover, even if A observes the execution of C and knows its type, A cannot infer specific details about the changes made to PT (computation confidentiality).*

HEOV is instantiated with two specific cryptographic primitives: (1) the BFV [9, 19] scheme, which achieves IND-CPA security based on the Decisional Ring Learning With Errors (D-RLWE) problem, and (2) Groth16 [26], which is a computationally sound and perfectly NIZKP system. When a probabilistic polynomial time attacker A corrupts a participating organization and observes all transactions, HEOV ensures *storage confidentiality*. This is achieved because A is computationally unable to distinguish the plaintext of a given ciphertext without the corresponding private key (decryption key), thanks to the IND-CPA security of the BFV scheme. Furthermore, throughout the transaction execution, transaction arguments are always concealed in ciphertext form. This ensures

that any two transactions referring to the same set of keys are probabilistically indistinguishable. In other words, *A* cannot gain any information beyond the identities of relevant keys during transaction execution, thereby achieving *computation confidentiality*.

6 Implementation

We implement HEOV on the codebase of HLF v2.2 [3]. The smart contract syntax of HEOV remains the same as HLF, but smart contracts need to be rewritten to use the libHEOV API for leveraging HE and NIZKP support, as discussed in §6.1.

6.1 libHEOV

libHEOV is a developer-friendly toolkit that minimizes the need for cryptography expertise and simplifies the use of HE and NIZKP functionalities. It harmonizes two mature libraries, lattigo [34] and gnark [13], which handle the arithmetic computations for HE and NIZKP operations, respectively.

Cryptography Arguments. libHEOV provides pre-defined cryptography arguments out-of-the-box (which will be discussed shortly) that balance security, speed, and memory consumption. libHEOV also offers high configurability, enabling developers to customize these arguments to achieve their desired trade-off between security and performance.

Homomorphic Encryption. libHEOV employs the BFV scheme [9, 19] as its FHE solution, ensuring 128-bit security in the default setting and supporting 58-bit unsigned integer arithmetic. HE multiplications in this setting take only tens of milliseconds on modern commodity servers [34]. Furthermore, the 58-bit unsigned integer is sufficient for many general applications without concerns about overflow.

Non-Interactive Zero-Knowledge Proof. libHEOV uses Groth16 [26] on the BN254 elliptic curve as its NIZKP solution. Like other systems relying on Groth16 zk-SNARKs [4, 46], our implementation requires a circuit-specific trusted setup, which can be performed using secure multi-party computation [5]. The overhead of the trusted setup process is considered trivial, as it is a one-shot procedure and therefore not included in the later evaluation section.

6.2 Rewrite Smart Contract With libHEOV

To enable HE and NIZKP support, developers need to manually rewrite a HLF smart contract according to the specifications stated below. Automated transpiling is not applicable here as it is challenging due to the potential ambiguity of smart contract semantics.

Use encrypted data type for private data. An encrypted variable is the ciphertext encrypted from its underlying plaintext. Throughout the execution of HE computation, the plaintext of it is never revealed without access to the relevant private key, ensuring the confidentiality of private data.

Use libHEOV functions to express HE and NIZKP operations. Replace standard addition and multiplication logic with the provided *Add* and *Mul* functions for HE computation. For NIZKP operations, use the *ProofGen* and *ProofVerify* functions to generate new proof instances and verify proofs with specific public input.

Use NIZKP instead of if-statements. Although most modern programming languages support if-statements or similar structures to control program flow based on conditional expressions, these statements cannot handle ciphertext directly due to the lack of support for direct comparison. As an alternative, HEOV employs NIZKP to simulate if-statements on the ciphertext. However, note that NIZKP cannot perfectly replace if-statements because, if a proof is verified as false, the program flow will exit without executing the remaining instructions.

7 Evaluation

We ran all experiments in a cluster with 20 machines, each equipped with a 2.60GHz E5-2690 CPU, 64GB memory, and a 40Gbps NIC. Each node and client was run in a separate docker container. A multi-host Docker Swarm network orchestrated all participant containers. The average node-to-node RTT was about 0.2 ms.

Baselines. We compared HEOV with ZeeStar (EOV), HEOV (w/o. correlated-merging), and HLF [3]. ZeeStar is an OE-based blockchain on the Ethereum platform, differing from the systems in our evaluation. For a fair comparison, we implemented ZeeStar (EOV) on the codebase of HLF and used it as a baseline instead of directly benchmarking ZeeStar [46]. ZeeStar (EOV) uses the exponential ElGamal encryption [33] as its additive homomorphic encryption solution and Groth16 as its NIZKP solution. We also developed HEOV (w/o. correlated-merging) representing HEOV without the correlated-merging protocol. Additionally, we evaluated HLF as it is the most popular permissioned blockchain framework in industry and academia [25, 40, 41].

Workloads. We evaluated four blockchains (HEOV, ZeeStar (EOV), HEOV (w/o. correlated-merging), and HLF) using two workloads.

The first workload was *SmallBank* [28], a widely used benchmark [40, 43] for evaluating blockchain systems that simulates bank business logic. The SmallBank workload was used to measure end-to-end performance (§7.1) and performance under no-resubmission attacks (§7.2). Due to the disparity in capabilities, we created three semantically equivalent implementations of SmallBank for each blockchain platform: an unencrypted version for HLF, an implementation using additive homomorphic encryption for ZeeStar (EOV), and an implementation using FHE with the libHEOV APIs (§6.1) for HEOV and HEOV (w/o. correlated-merging). Each experiment started with ten organizations, each with 100 ac-

Figure 4: End-to-end performance of HEOV and HLF under applications with different conflict ratios.

counts initialized with the same balance. Subsequently, we generated a set of *transfer* transactions, where a randomly selected account $Acct_1$ transferred a specific number of tokens to another randomly chosen account $Acct_2$. Note that $Acct_1$ and $Acct_2$ may or may not belong to the same organization.

Our second workload is a microbenchmark designed to showcase the performance improvement introduced by the correlated-merging protocol. It comprises two smart contracts, SC_A and SC_M, which execute homomorphic additive and multiplicative transactions, respectively. The ratio between these two transaction types is configurable to simulate various scenarios. We assessed this workload exclusively on HEOV and HEOV (w/o. correlated-merging) (§7.3), as HLF does not support on-ciphertext computation and ZeeStar (EOV) does not utilize the correlated-merging protocol, therefore not applicable to this evaluation.

Metrics. We evaluated two metrics: *effective throughput*, which represents the average number of valid client transactions committed to the blockchain per second, excluding conflicting aborted transactions, and *commit latency* (also known as *end-to-end latency*), which measures the duration from client submission to transaction commitment. Note that if a transaction fails to commit, HEOV's lead executor will automatically resubmit the transaction until it is eventually committed. Additionally, we reported the 99th percentile commit latency and the cumulative distribution function (CDF) of the transactions' commit latency.

Evaluation methodology. We developed a distributed benchmark using Tape [29], an efficient benchmark tool for HLF. The benchmark spawned 128 clients across multiple servers to prevent the benchmarking tool from becoming a bottleneck.

Each system was configured with ten executors and four orderers, utilizing BFT-SMaRT [7] as the consensus protocol. The default block size was set to 100 transactions, a commonly used setting in relevant studies [41, 43]. To mimic real-world scenarios, we designated 1% of the accounts as *Hot Accounts*, and the *Conflict Ratio* represented the probability of each transaction accessing these hot accounts, with a default value of 50%. For each evaluation, we performed ten runs and reported the average values for each metric. Our evaluation focused on three primary questions.

§7.1 How efficient is HEOV in a benign environment?

§7.2 How robust is HEOV against no-resubmission attacks?

§7.3 How efficient is HEOV across diverse applications?

7.1 End-to-End Performance

We first conducted experiments in benign environments where the network was stable and all participants behaved correctly. All systems are evaluated with different conflict ratios of 10%, 50%, and 90%, respectively.

HEOV outperformed both ZeeStar (EOV) and HEOV (w/o. correlated-merging) in terms of effective throughput, achieving increases up to 7.14×. In Figure 4, HEOV achieved throughput values of 255 TPS, 307 TPS, and 343 TPS at conflict ratios of 10%, 50%, and 90%, respectively. In contrast, HEOV (w/o. correlated-merging) only achieved 235 TPS, 95 TPS, and 55 TPS, demonstrating a significant performance gap between HEOV and HEOV (w/o. correlated-merging), especially with higher conflict ratios. ZeeStar (EOV) performed even worse, achieving only 213 TPS, 85 TPS, and 48 TPS.

The average end-to-end latency for HEOV is 2.7s, which is only 87% of the baseline ZeeStar (EOV) (3.1s) and slightly higher than HEOV (w/o. correlated-merging) (2.4s). The additional latency in HEOV is due to the pending time required for merging transactions, as introduced by the correlated-merging protocol (§4.2). However, the latency overhead in HEOV was merely 10%, while the throughput gain could reach up to 5.2×. Even when using FHE, which incurs higher costs than the PHE solution of ZeeStar (EOV), HEOV demonstrated shorter average end-to-end latency compared to ZeeStar (EOV). This advantage is due to HEOV's lightweight NIZKP (§4.2), eliminating the need to prove the correctness of individual data updates in each transaction, a requirement in ZeeStar (EOV).

HEOV is particularly suitable for applications with conflicting transactions thanks to its correlated-merging protocol (§4.2). Surprisingly, in Figure 6, HEOV significantly outperformed HLF with 1.54x higher throughput at a 90% conflict ratio, despite the extra overhead from HE and NIZKP gen-

System	Average Latency (s)					99% tail latency (s)
	P	E	O	V	E2E	
HEOV	0.75	0.85	0.89	0.23	2.72	3.28
HEOV (w/o. correlated-merging)	0.47	0.87	0.88	0.24	2.46	7.46
ZeeStar (EOV)	0.26	1.66	0.87	0.32	3.11	6.85
HLF	0.03	0.41	0.91	0.19	1.54	5.03

Table 2: Latency of systems under test in Figure 4b and 4d. The "P" phase represents the preparation period before execution.

(a) Effective throughput   (b) Latency distribution during no-resubmission attack

Figure 5: Performance under no-resubmission attack

eration. In less conflicting scenarios, HEOV was not as performant as HLF, but it still achieved higher throughput than ZeeStar (EOV) and HEOV (w/o. correlated-merging) in all three settings, albeit at the cost of ensuring data confidentiality.

As shown in Figure 4, HEOV exhibited an increasing trend in throughput as the conflict ratio increased, performing even better in less conflicting scenarios. In contrast, ZeeStar (EOV), HEOV (w/o. correlated-merging), and HLF suffered from a significant drop in throughput. This is because HEOV's correlated-merging protocol effectively reduces a substantial portion of conflicts, resulting in fewer cryptographic operations and conflicting aborts. Unlike the other systems, HEOV efficiently handles conflicts, allowing for sustained high throughput.

Figure 4d shows that HEOV's average latency (2.7s) was higher than that of HLF (1.6s), due to the extra overhead introduced by the HE and NIZKP operations. However, the increase in latency is fully justified by HEOV's confidentiality guarantee for general smart contracts. Furthermore, it is noteworthy that HEOV achieved a shorter 99%-tail latency compared to HLF, and all transaction latencies were concentrated within a narrow range. This indicates that the correlated-merging protocol employed by HEOV generally reduces the need for transaction re-submission.

Table 2 provides a more detailed insight into the latency for each phase: preparation (**P**), execution (**E**), ordering (**O**), and validation (**V**). HEOV incurred extra overhead primarily in the preparation and execution phases due to encryption, merging, expensive operations like HE evaluation, NIZKP generation, and verification. The validation phase also experienced a slight latency increase due to NIZKP verification. The ordering phase, on the other hand, remains unaffected as expected. These results highlight the trade-off between the confidentiality guarantee provided by HEOV and the extra overhead from FHE and NIZKP operations.

Overall, HEOV offers a compelling combination of high performance and strong security guarantees, making it well-suited for applications that prioritize data privacy and involve conflicting transactions.

7.2 Robustness to No-Resubmission Attacks

We conducted no-resubmission attacks on four systems (Figure 5) involving four benign organizations (O_B) and one malicious organization (O_M). A victim organization O_V was selected from O_B. O_M created transactions involving O_V and intentionally did not resubmit them in order to attack O_V. The experiments were divided into three consecutive periods: *Period I: pre-attack*, *Period II: attack*, and *Period III: post-attack*. Metrics are shown in Figure 5.

Figure 5a shows that initially, all four systems experienced varying levels of performance degradation during the no-resubmission attack. However, HEOV and HEOV (w/o. correlated-merging)'s throughput quickly resumed to the *Period I* level, while ZeeStar (EOV) and HLF remained at a low level. This showcases the effectiveness of the countermeasure, where executors of benign organizations block malicious lead executors, preventing monopolization and ensuring fair access.

In Figure 5b, both HEOV and HEOV (w/o. correlated-merging) exhibited much better average latency than ZeeStar (EOV) and HLF during the no-resubmission attack. They efficiently detected and rejected malicious attackers, preserving computation power for benign clients and mitigating latency impact.

Overall, the results in Figure 5 emphasize that HEOV is particularly suitable for privacy-sensitive applications involving malicious participants. It incorporates countermeasures against no-resubmission attacks, ensuring fair access, and reducing latency.

7.3 Performance under Diverse Applications

The *SmallBank* workload represents real-world applications like token transfers, focusing on homomorphic addition rather than multiplication. To fully benchmark the performance improvement brought by correlated-merging on both addition and multiplication, we developed the second workload introduced in §7, and conducted three experiments to evaluate the performance of HEOV and HEOV (w/o. correlated-merging): (1) (**Add**) all transactions only invoked SC_A; (2) (**Mul**) all

transactions only invoked SC_M; and (3) (**Add-Mul**) 50% of transactions invoked SC_A, while the other 50% invoked SC_M.

Results in Figure 6 show that HEOV's correlated-merging protocol supports both addition and multiplication operations, making it ideal for applications that require one or both of these types of operations, such as cryptocurrency [17] and finance [49]. Figure 6 also demonstrates that HEOV significantly outperformed HEOV (w/o. correlated-merging) in all three experiments, with throughput enhancements of 3.2x, 5.4x, and 7.8x for addition-only, multiplication-only, and hybrid applications, respectively. These results highlight the effectiveness of HEOV's correlated-merging protocol, particularly when both addition and multiplication are involved. Additionally, the results suggest that the correlated-merging protocol excels when multiplication is the major HE computation of the business logic.

In summary, HEOV's correlated-merging protocol has delivered substantial performance advantages to HEOV without altering the transaction semantics. This breakthrough enables HEOV to safeguard the privacy of sensitive user data with cryptographic guarantees while maintaining high performance.

7.4 Lessons Learned

HEOV has two limitations. First, HEOV is designed specifically for blockchain applications involving conflicting transactions, such as token transfers [17]. It employs the correlated-merging protocol that effectively reduces the rate of conflicting aborts and minimizes the computational overhead of expensive HE operations. However, for conflict-free applications, HEOV performs similarly to the baseline. Note that conflicts are common in typical real-world blockchain applications, and achieving conflict-free applications would require significant modifications to the application logic, such as rewriting the smart contract using CRDT [36]. This may not be feasible for widely used blockchain applications, such as financial trading [36].

Second, the current implementation of HEOV only supports HE computations on bounded unsigned integers due to its adoption of the BFV scheme [9, 19]. However, incorporating more capable FHE schemes like CKKS [12]) would enable HEOV to easily support HE computations on floating-point numbers. This would broaden the range of applications that HEOV can accommodate, enhance its flexibility, and improve its usability for real-world blockchain applications.

8 Conclusion

We present HEOV, the first privacy-preserving EOV permissioned blockchain that efficiently achieves both BFT safety and data confidentiality. HEOV harmonizes the best of fully homomorphic encryption, non-interactive zero-knowledge proof, and the EOV workflow, while circumventing their


Figure 6: HEOV's optimization improvement under different combinations of HE computation.

drawbacks via a novel correlated-merging homomorphic encryption protocol. Extensive evaluation results on typical blockchain applications with diverse computations show that HEOV is general, highly efficient, and secure compared to the baselines. HEOV is open-sourced, and its code is released on https://github.com/asplos24p5/heov.

References

[1] Ittai Abraham, Guy Gueta, and Dahlia Malkhi. Hot-stuff the linear, optimal-resilience, one-message BFT devil, 2018.

[2] Farhana Aleen and Nathan Clark. Commutativity analysis for software parallelization: letting program transformations see the big picture. *ACM Sigplan Notices*, 44(3):241–252, 2009.

[3] Elli Androulaki, Artem Barger, Vita Bortnikov, Christian Cachin, Konstantinos Christidis, Angelo De Caro, David Enyeart, Christopher Ferris, Gennady Laventman, Yacov Manevich, et al. Hyperledger fabric: a distributed operating system for permissioned blockchains. In *Proceedings of the thirteenth EuroSys conference*, pages 1–15, New York, NY, USA, 2018. ACM.

[4] Nick Baumann, Samuel Steffen, Benjamin Bichsel, Petar Tsankov, and Martin T. Vechev. zkay v0.2: Practical data privacy for smart contracts, 2020.

[5] Eli Ben-Sasson, Alessandro Chiesa, Matthew Green, Eran Tromer, and Madars Virza. Secure sampling of public parameters for succinct zero knowledge proofs. In *2015 IEEE Symposium on Security and Privacy*, pages 287–304, 2015.

[6] Alysson Bessani, João Sousa, and Eduardo EP Alchieri. State machine replication for the masses with bft-smart. In *2014 44th Annual IEEE/IFIP International Conference on Dependable Systems and Networks*, pages 355–362, 1730 Massachusetts Ave., NW Washington, DCUnited States, 2014. IEEE, IEEE Computer Society.

[7] Alysson Bessani, João Sousa, and Eduardo E.P. Alchieri. State machine replication for the masses with bft-smart.

In *2014 44th Annual IEEE/IFIP International Conference on Dependable Systems and Networks*, pages 355–362, 2014.

[8] Nir Bitansky, Ran Canetti, Alessandro Chiesa, and Eran Tromer. From extractable collision resistance to succinct non-interactive arguments of knowledge, and back again. In *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, pages 326–349, New York, NY, USA, 2012. ACM.

[9] Zvika Brakerski. Fully homomorphic encryption without modulus switching from classical gapsvp. In *Advances in Cryptology–CRYPTO 2012: 32nd Annual Cryptology Conference, Santa Barbara, CA, USA, August 19-23, 2012. Proceedings*, pages 868–886, Heidelberg, 2012. Springer, Springer Berlin.

[10] Zvika Brakerski, Craig Gentry, and Vinod Vaikuntanathan. (leveled) fully homomorphic encryption without bootstrapping. *ACM Transactions on Computation Theory (TOCT)*, 6(3):1–36, 2014.

[11] Miguel Castro, Barbara Liskov, et al. Practical byzantine fault tolerance. In *OSDI*, volume 99, pages 173–186, New York, NY, USA, 1999. ACM.

[12] Jung Hee Cheon, Andrey Kim, Miran Kim, and Yongsoo Song. Homomorphic encryption for arithmetic of approximate numbers. In *Advances in Cryptology–ASIACRYPT 2017: 23rd International Conference on the Theory and Applications of Cryptology and Information Security, Hong Kong, China, December 3-7, 2017, Proceedings, Part I 23*, pages 409–437, Hong Kong, China, 2017. Springer, Springer.

[13] consensys. Gnark, 2023.

[14] Cosmos. validators overview, 2023.

[15] Thi Van Thao Doan, Mohamed-Lamine Messai, Gérald Gavin, and Jérôme Darmont. A survey on implementations of homomorphic encryption schemes. *The Journal of Supercomputing*, pages 1–42, 2023.

[16] Jacob Eberhardt and Stefan Tai. Zokrates-scalable privacy-preserving off-chain computations. In *2018 IEEE International Conference on Internet of Things (iThings) and IEEE Green Computing and Communications (GreenCom) and IEEE Cyber, Physical and Social Computing (CPSCom) and IEEE Smart Data (SmartData)*, pages 1084–1091, Halifax, Nova Scotia, Canada, 2018. IEEE, IEEE.

[17] Ethreum. Erc-20 token standard, 2023.

[18] HyperLedger Fabric. Case Study:How Walmart brought unprecedented transparency to the food supply chain with Hyperledger Fabric. https://www.hyperledger.org/learn/publications/walmart-case-study, 2022.

[19] Junfeng Fan and Frederik Vercauteren. Somewhat practical fully homomorphic encryption. Cryptology ePrint Archive, Paper 2012/144, 2012. https://eprint.iacr.org/2012/144.

[20] Uriel Feige, Dror Lapidot, and Adi Shamir. Multiple noninteractive zero knowledge proofs under general assumptions. *SIAM Journal on computing*, 29(1):1–28, 1999.

[21] Caroline Fontaine and Fabien Galand. A survey of homomorphic encryption for nonspecialists. *EURASIP Journal on Information Security*, 2007:1–10, 2007.

[22] Yossi Gilad, Rotem Hemo, Silvio Micali, Georgios Vlachos, and Nickolai Zeldovich. Algorand: Scaling byzantine agreements for cryptocurrencies. In *Proceedings of the 26th Symposium on Operating Systems Principles*, SOSP '17, page 51–68, New York, NY, USA, 2017. Association for Computing Machinery.

[23] Guy Golan Gueta, Ittai Abraham, Shelly Grossman, Dahlia Malkhi, Benny Pinkas, Michael Reiter, Dragos-Adrian Seredinschi, Orr Tamir, and Alin Tomescu. Sbft: A scalable and decentralized trust infrastructure. In *2019 49th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, pages 568–580, Portland, OR, USA, 2019. IEEE.

[24] Google. Build simple, secure, scalable systems with go.

[25] Christian Gorenflo, Stephen Lee, Lukasz Golab, and Srinivasan Keshav. Fastfabric: Scaling hyperledger fabric to 20 000 transactions per second. *International Journal of Network Management*, 30(5):e2099, 2020.

[26] Jens Groth. On the size of pairing-based non-interactive arguments. In Marc Fischlin and Jean-Sébastien Coron, editors, *Advances in Cryptology – EUROCRYPT 2016*, pages 305–326, Berlin, Heidelberg, 2016. Springer Berlin Heidelberg.

[27] Jens Groth and Mary Maller. Snarky signatures: Minimal signatures of knowledge from simulation-extractable snarks. In Jonathan Katz and Hovav Shacham, editors, *Advances in Cryptology – CRYPTO 2017*, pages 581–612, Cham, 2017. Springer International Publishing.

[28] H-Store. H-store: Smallbank benchmark, 2013.

[29] Hyperledger-TWGC. Hyperledger-twgc/tape: A simple traffic generator for hyperledger fabric, 2023.

[30] Hui Kang, Ting Dai, Nerla Jean-Louis, Shu Tao, and Xi-aohui Gu. Fabzk: Supporting privacy-preserving, auditable smart contracts in hyperledger fabric. In *2019 49th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN)*, pages 543–555, Portland, Oregon, USA, 2019. IEEE, IEEE.

[31] Yang Liu, Debiao He, Mohammad S Obaidat, Neeraj Kumar, Muhammad Khurram Khan, and Kim-Kwang Raymond Choo. Blockchain-based identity management systems: A review. *Journal of network and computer applications*, 166:102731, 2020.

[32] Medicalchain. Medicalchain. https://medicalchain.com, 2022.

[33] Andreas V Meier. The elgamal cryptosystem. In *Joint Advanced Students Seminar*, 2005.

[34] Christian Vincent Mouchet, Jean-Philippe Bossuat, Juan Ramón Troncoso-Pastoriza, and Jean-Pierre Hubaux. Lattigo: A multiparty homomorphic encryption library in go. In *Proceedings of the 8th Workshop on Encrypted Computing and Applied Homomorphic Cryptography*, pages 6. 64–70, ., 2020. HomomorphicEncryption.org.

[35] Satoshi Nakamoto. Bitcoin: A peer-to-peer electronic cash system, 2008.

[36] Pezhman Nasirifard, Ruben Mayer, and Hans-Arno Jacobsen. Fabriccrdt: A conflict-free replicated datatypes approach to permissioned blockchains. In *Proceedings of the 20th International Middleware Conference*, pages 110–122, 2019.

[37] NCCGroup, 2016.

[38] Pascal Paillier. Public-key cryptosystems based on composite degree residuosity classes. In Jacques Stern, editor, *Advances in Cryptology — EUROCRYPT '99*, pages 223–238, Berlin, Heidelberg, 1999. Springer Berlin Heidelberg.

[39] Pascal Paillier and David Pointcheval. Efficient public-key cryptosystems provably secure against active adversaries. In Kwok-Yan Lam, Eiji Okamoto, and Chaoping Xing, editors, *Advances in Cryptology - ASIACRYPT'99*, pages 165–179, Berlin, Heidelberg, 1999. Springer Berlin Heidelberg.

[40] Ji Qi, Xusheng Chen, Yunpeng Jiang, Jianyu Jiang, Tianxiang Shen, Shixiong Zhao, Sen Wang, Gong Zhang, Li Chen, Man Ho Au, and Heming Cui. Bidl: A high-throughput, low-latency permissioned blockchain framework for datacenter networks. In *Proceedings of the ACM SIGOPS 28th Symposium on Operating Systems Principles*, SOSP '21, page 18–34, New York, NY, USA, 2021. Association for Computing Machinery.

[41] Pingcheng Ruan, Dumitrel Loghin, Quang-Trung Ta, Meihui Zhang, Gang Chen, and Beng Chin Ooi. A transactional perspective on execute-order-validate blockchains. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*, SIGMOD '20, page 543–557, New York, NY, USA, 2020. Association for Computing Machinery.

[42] Microsoft SEAL (release 4.0). https://github.com/Microsoft/SEAL, March 2022. Microsoft Research, Redmond, WA.

[43] Ankur Sharma, Felix Martin Schuhknecht, Divya Agrawal, and Jens Dittrich. Blurring the lines between blockchains and database systems: The case of hyperledger fabric. In *Proceedings of the 2019 International Conference on Management of Data*, SIGMOD '19, page 105–122, New York, NY, USA, 2019. Association for Computing Machinery.

[44] Ravital Solomon, Rick Weber, and Ghada Almashaqbeh. smartfhe: Privacy-preserving smart contracts from fully homomorphic encryption. *Cryptology ePrint Archive*, 2021.

[45] Chrysoula Stathakopoulou, Tudor David, and Marko Vukolic. Mir-bft: High-throughput BFT for blockchains, 2019.

[46] Samuel Steffen, Benjamin Bichsel, Roger Baumgartner, and Martin Vechev. Zeestar: Private smart contracts by homomorphic encryption and zero-knowledge proofs. In *2022 IEEE Symposium on Security and Privacy (SP)*, pages 179–197, San Francisco, California, USA, 2022. IEEE.

[47] Samuel Steffen, Benjamin Bichsel, Mario Gersbach, Noa Melchior, Petar Tsankov, and Martin Vechev. Zkay: Specifying and enforcing data privacy in smart contracts. In *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*, CCS '19, page 1759–1776, New York, NY, USA, 2019. Association for Computing Machinery.

[48] Samuel Steffen, Benjamin Bichsel, and Martin Vechev. Zapper: Smart contracts with data and identity privacy. In *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security*, pages 2735–2749, 2022.

[49] taXchain. taxchain provides a faster, better, cheaper way to complete eu tax forms using hyperledger fabric, 2023.

[50] Alexander Thomson, Thaddeus Diamond, Shu-Chun Weng, Kun Ren, Philip Shao, and Daniel J. Abadi. Calvin: Fast distributed transactions for partitioned database systems. In *Proceedings of the 2012 ACM*

SIGMOD International Conference on Management of Data, SIGMOD '12, page 1–12, New York, NY, USA, 2012. Association for Computing Machinery.

[51] Alexander Viand, Patrick Jattke, and Anwar Hithnawi. Sok: Fully homomorphic encryption compilers. In *2021 IEEE Symposium on Security and Privacy (SP)*, pages 1092–1108, San Francisco, CA, USA, 2021. IEEE.

[52] Paul Voigt and Axel Von dem Bussche. The eu general data protection regulation (gdpr). *A Practical Guide, 1st Ed., Cham: Springer International Publishing*, 10(3152676):10–5555, 2017.

[53] Gavin Wood et al. Ethereum: A secure decentralised generalised transaction ledger. *Ethereum project yellow paper*, 151(2014):1–32, 2014.

[54] Kimchai Yeow, Abdullah Gani, Raja Wasim Ahmad, Joel JPC Rodrigues, and Kwangman Ko. Decentralized consensus for edge-centric internet of things: A review, taxonomy, and research issues. *IEEE Access*, 6:1513–1524, 2017.

[55] zcash. What are zk-snarks?, Jun 2022.